

```
Activities Terminal Oct 6 10:20 hadoop@ab4cseibm133: ~
hadoop@ab4cseibm133:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [ab4cseibm133]
Starting resourcemanager
Starting nodemanagers
hadoop@ab4cseibm133:~$ jps
22513 DataNode
22996 ResourceManager
23238 NodeManager
23591 Jps
22364 NameNode
22717 SecondaryNameNode
hadoop@ab4cseibm133:~$ hadoop fs -mkdir hdfs://localhost:9000/ex10
mkdir: 'hdfs://localhost:9000/ex10': File exists
hadoop@ab4cseibm133:~$ hadoop fs -put '/home/hadoop/student_details.txt' /ex10
hadoop@ab4cseibm133:~$ jps
2022-10-06 10:05:13,220 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL
2022-10-06 10:05:13,220 INFO pig.ExecTypeProvider: Picked LOCAL as the ExecType
2022-10-06 10:05:13,277 [main] INFO org.apache.pig.Main - Apache Pig version 0.17.0 (r1797386) compiled Jun 02 2017, 15:41:58
2022-10-06 10:05:13,277 [main] INFO org.apache.pig.Main - Logging error messages to: /home/hadoop/pig_1665030913276.log
2022-10-06 10:05:13,328 [main] INFO org.apache.pig.impl.util.Utils - Default bootstrap file /home/hadoop/pigbootstrap not found
2022-10-06 10:05:13,441 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapred
ce.jobtracker.address
2022-10-06 10:05:13,443 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to hadoop file system at: fl
e:///
2022-10-06 10:05:13,502 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.
bytes-per-checksum
2022-10-06 10:05:13,514 [main] INFO org.apache.pig.PigServer - Pig Script ID for the session: PIG-default-192573b8-a43f-465a-b4a8-8e07b0d819f
b
2022-10-06 10:05:13,514 [main] WARN org.apache.pig.PigServer - ATS is disabled since yarn.timeline.service.enabled set to false
grunt> student_details - LOAD 'hdfs://localhost:9000/pig_data/student_details.txt' USING PigStorage(',') as (id:int, firstname:chararray, last
name:chararray, age:int, phone:chararray, city:chararray);
2022-10-06 10:06:14,346 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.
bytes-per-checksum
grunt>
```

```
Activities Terminal Oct 6 10:18 hadoop@ab4cseibm133: ~
Counters:
Total records written : 8
Total bytes written : 0
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_local1830704608_0003 -> job_local773754855_0004,
job_local773754855_0004 -> job_local2015787209_0005,
job_local2015787209_0005

2022-10-06 10:13:52,689 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,691 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,692 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,696 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,698 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,699 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,703 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,705 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,706 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:13:52,708 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce.Launcher - Success!
2022-10-06 10:13:52,709 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.
bytes-per-checksum
2022-10-06 10:13:52,709 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized
2022-10-06 10:13:52,711 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input files to process : 1
2022-10-06 10:13:52,711 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths to process : 1
(0, Bharathi, Namblayar, 24, 9848022333, Chennai)
(7, Konal, Nayak, 24, 9848022334, trivendram)
(6, Archana, Mishra, 23, 9848022335, Chennai)
(5, Trupthi, Mohanthy, 23, 9848022336, Bhuvaneshwar)
(3, Rajesh, Khanna, 22, 9848022339, Delhi)
(2, Siddarth, Battacharya, 22, 9848022338, Kolkata)
(4, Preethi, Agarwal, 21, 9848022330, Pune)
(1, Rajiv, Reddy, 21, 9848022337, Hyderabad)
grunt>
```

```
Activities Terminal Oct 6 10:35 hadoop@ab4cseibm133: ~
Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime Median
ReduceTime Alias Feature Outputs
job_local664789670_0006 1 1 n/a n/a n/a n/a n/a n/a group_data,student_details GROUP_
BY file:/tmp/temp1257080311/tmp-1137725553,

Input(s):
Successfully read 8 records (1356 bytes) from: "hdfs://localhost:9000/ex10/student_details.txt"

Output(s):
Successfully stored 4 records in: "file:/tmp/temp1257080311/tmp-1137725553"

Counters:
Total records written : 4
Total bytes written : 0
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_local664789670_0006

2022-10-06 10:35:26,077 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:35:26,078 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:35:26,080 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!
2022-10-06 10:35:26,082 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce.Launcher - Success!
2022-10-06 10:35:26,083 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.
bytes-per-checksum
2022-10-06 10:35:26,083 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized
2022-10-06 10:35:26,083 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input files to process : 1
2022-10-06 10:35:26,085 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths to process : 1
(21, ((4, Preethi, Agarwal, 21, 9848022330, Pune), (1, Rajiv, Reddy, 21, 9848022337, Hyderabad)))
(22, ((3, Rajesh, Khanna, 22, 9848022339, Delhi), (2, Siddarth, Battacharya, 22, 9848022338, Kolkata)))
(23, ((6, Archana, Mishra, 23, 9848022335, Chennai), (5, Trupthi, Mohanthy, 23, 9848022336, Bhuvaneshwar)))
(24, ((8, Bharathi, Namblayar, 24, 9848022333, Chennai), (7, Konal, Nayak, 24, 9848022334, trivendram)))
grunt>
```


Join:

```
customers = LOAD 'hdfs://localhost:9000/ex10/customers.txt' USING PigStorage(',') as (id:int,  
name:chararray, age:int, address:chararray, salary:int);
```

```
orders = LOAD 'hdfs://localhost:9000/ex10/orders.txt' USING PigStorage(',') as (oid:int,  
date:chararray, customer_id:int, amount:int);
```

```
customers1 = LOAD 'hdfs://localhost:9000/ex10/customers.txt' USING PigStorage(',') as (id:int,  
name:chararray, age:int, address:chararray, salary:int);
```

```
customers2 = LOAD 'hdfs://localhost:9000/ex10/customers.txt' USING PigStorage(',') as (id:int,  
name:chararray, age:int, address:chararray, salary:int);
```

Group:

```
group_data = GROUP student_details by age;  
Dump group_data;
```

Project:

```
foreach_data = FOREACH student_details GENERATE id,age,city;  
Dump foreach_data;
```

Filter:

```
filter_data = FILTER student_details BY city == 'Chennai';  
Dump filter_data;
```

Sort:

```
order_by_data = ORDER student_details BY age DESC;  
Dump order_by_data;
```