

#Commands

```
create database DBname;
use dbname;
show databases;
drop database Dbname;
show databases;
create table TBname (columnname datatype, columnname2 datatype);
describe Tbname;
load data local inpath '/home/codegyani/hive/emp_details' into table TBname;
(or)
insert into Tbname values(value1,value2);
Select * from Tbname;
drop table TBname;
show tables;
Alter table old_table_name rename to new_table_name;
show tables;
Alter table table_name add columns(column_name datatype);
Alter table table_name change old_column_name new_column_name datatype;
describe tbname;
alter table TBname replace columns( columnname datatype, col2name datatype);
describe TBname;
CREATE VIEW view_name AS SELECT ...;
DROP VIEW view_name ;
```

Output:

```
Oct 27 12:23
hadoop@abi-cse108: ~/apache-hive-3.1.2-bin
... 18 more
hadoop@abi-cse108:~$ cd $HIVE_HOME
hadoop@abi-cse108:~/apache-hive-3.1.2-bin$ bin/hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.3/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = ca23c999-c087-496b-a669-c9be9fbc0fe2

Logging initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async
true
Hive Session ID = e9069e3f-4b04-40c8-9e43-d2f52d9fbed0
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> create database DB;
OK
Time taken: 0.787 seconds
hive> use DB;
OK
Time taken: 0.021 seconds
hive> show databases;
OK
bighive
data
db
default
g)
student
Time taken: 0.167 seconds, Fetched: 6 row(s)
hive> create table list(id int,veg int);
OK
Time taken: 0.88 seconds
hive> describe list;
OK
id          int
veg         int
Time taken: 0.225 seconds, Fetched: 2 row(s)
```

```
Oct 27 12:24
hadoop@abi-cse108: ~/apache-hive-3.1.2-bin
FAILED: ParseException line 1:7 cannot recognize input near 'insert' 'int' 'list' in insert clause
Hive> insert into list values(1,45);
Query ID = hadoop_20221027120311_1a3d5ca4-2808-46c0-a902-ed4237748530
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Starting Job = job_1666851099333_0001, Tracking URL = http://abi-cse108:8088/proxy/application_1666851099333_0001/
Kill Command = /home/hadoop/hadoop-3.2.3/bin/mapred job -kill job_1666851099333_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-10-27 12:03:11,987 Stage-1 map = 0%, reduce = 0%
2022-10-27 12:03:48,964 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.23 sec
2022-10-27 12:03:57,123 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.98 sec
MapReduce Total cumulative CPU time: 3 seconds 980 msec
Ended Job = job_1666851099333_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/db.db/list/.hive-staging_hive_2022-10-27_12-03_13_932_1282054495814042905-1
/ext-10000
Loading data to table db.list
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.98 sec HDFS Read: 14193 HDFS Write: 229 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 980 msec
OK
Time taken: 48.906 seconds
hive> insert into list values(2,78);
Query ID = hadoop_20221027120449_451ef85e-4484-499b-93d0-14655d5e784c
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
```

```
Oct 27 12:25
hadoop@abi-cse108: ~/apache-hive-3.1.2-bin
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.9 sec HDFS Read: 14231 HDFS Write: 229 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 900 msec
OK
Time taken: 24.632 seconds
hive> select * from list;
OK
1      45
2      78
3      07
4      99
Time taken: 0.324 seconds, Fetched: 4 row(s)
hive> alter table list rename to marks;
OK
Time taken: 0.202 seconds
hive> describe marks;
OK
id          int
veg         int
Time taken: 0.080 seconds, Fetched: 2 row(s)
hive> alter table marks add columns(nonveg int);
OK
Time taken: 0.25 seconds
hive> describe marks;
OK
id          int
veg         int
nonveg      int
Time taken: 0.031 seconds, Fetched: 3 row(s)
hive> create view bob as select * from marks where veg>80;
OK
Time taken: 0.291 seconds
hive> select * from bob;
OK
3      87      NULL
4      99      NULL
Time taken: 0.196 seconds, Fetched: 2 row(s)
hive>
```